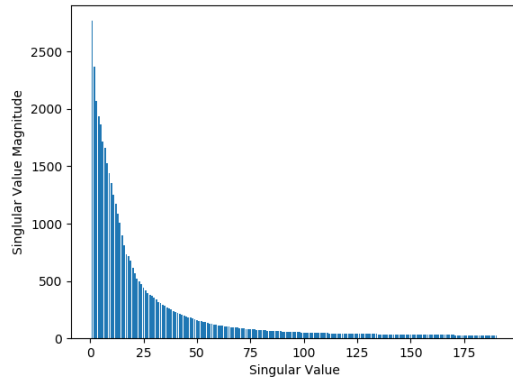
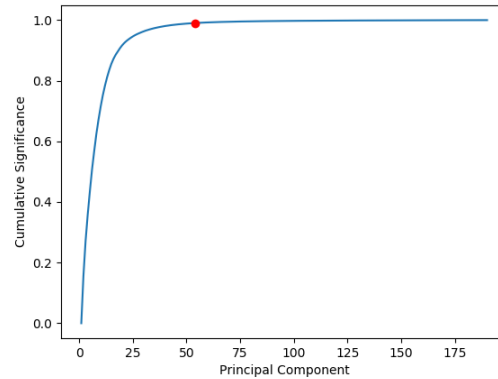


a)



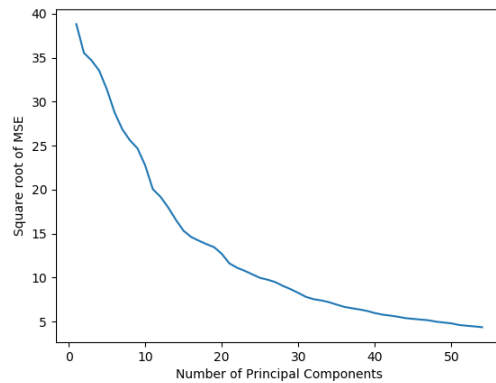
*Figure 1: Plot of singular values for corresponding principal components of 190/200 neutral faces*



*Figure 2: Plot of cumulative significance as additional principal components are used*

I decided to use the number of principal components that would encompass 99% of the total sum of the magnitudes of all the singular values. Since the significance of a principal component is determined by the magnitude of its singular value (larger singular values represent larger variance of the data in the direction of the corresponding eigenvector), it is reasonable to expect that if the sum of corresponding eigenvalues for the first 54 principal components equates to 99% of the total sum of all the eigenvalues – then these 54 principal components are able to encapsulate 99% of the data required to reconstruct an image. This allows me to decrease the number of principal components by 72% while only losing 1% of the data (assuming this heuristic approach equates to this conclusion). I have no proof of the validity of this approach – it is just an intuitive, heuristic approach obtained by observing Figures 1 and 2.

b)

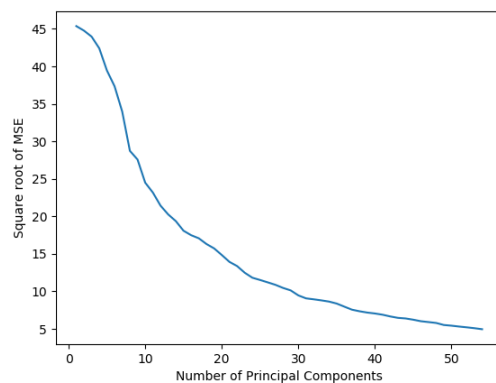


*Figure 3: Square root of MSE between neutral face 4a and its reconstruction vs number of principal components used for the reconstruction*

The square root of the mean squared error was plotted to aid in the visualization of the plot. As the number of principal components used is increased the MSE decreases exponentially and reaches a minimum value of approximately 25 ( $5^2$ ). This equates to approximately 1.95% of error\*, which is not far off from the 1% goal.

\* If the average difference between each pixel is 5 for pixels that range from 0 to 255, then the average percentage of difference between the images is calculated by  $\frac{5}{256} * 100 = 1.953125\%$ .

c)

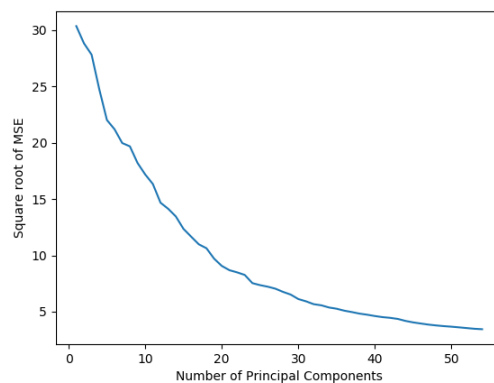


*Figure 4: Square root of MSE between smiling face 4b and its reconstruction vs number of principal components used for the reconstruction*

The square root of the mean squared error was plotted to aid in the visualization of the plot. As the number of principal components used is increased the MSE decreases exponentially and, just as Figure 3 does, reaches a minimum value of approximately 25 ( $5^2$ ). This equates to approximately 1.95% of error\*, which is not far off from the 1% goal. One noticeable difference is that the MSE hits a maximum value of approximately 45 in Figure 4, whereas the maximum value in Figure 3 is approximately 39. Also, the decrease in Figure 4 has a greater initial slope – indicating the greater impact of additional principal components on the MSE when principal component levels are below ~10. Since the principal components were created from neutral faces, it makes sense that the principal components would perform better at recreating a neutral face as opposed to a smiling one. This seems to be true when the number of principal components is less than or equal to approximately 10. After that, however, the principal components are just as successful in recreating a smiling face as they are in recreating a neutral face. However, this is only one example and it would be interesting to see if this trend holds for all pairs of neutral and smiling faces.

\* If the average difference between each pixel is 5 for pixels that range from 0 to 255, then the average percentage of difference between the images is calculated by  $\frac{5}{256} * 100 = 1.953125\%$ .

d)

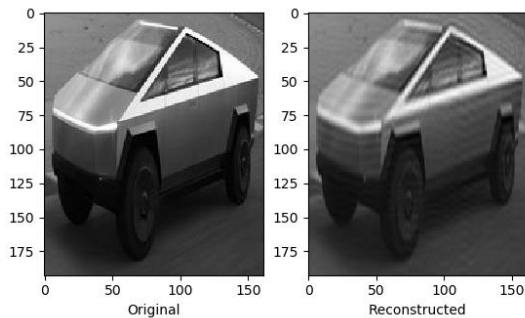


*Figure 5: Square root of MSE between neutral face 195a and its reconstruction vs number of principal components used for the reconstruction*

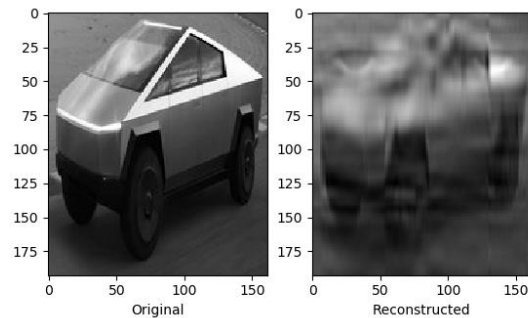
The square root of the mean squared error was plotted to aid in the visualization of the plot. As the number of principal components used is increased the MSE decreases exponentially and reaches a minimum value of approximately 16 ( $4^2$ ). This equates to approximately 1.56% of error\*, which is even closer to the 1% goal than previous attempts. One noticeable difference is that the MSE hits a maximum value of approximately 31 in Figure 5, whereas the maximum values were approximately 39 and 45 in Figures 3 and 4 respectively. It's interesting that the minimum MSE values are approximately the same, if not smaller, when the image to reconstruct is not used in the initial construction of the principal components. However, it may just have been that the faces selected for recreation in Figures 4 and 5 were easy to reconstruct for the set of principal components.

\* If the average difference between each pixel is 4 for pixels that range from 0 to 255, then the average percentage of difference between the images is calculated by  $\frac{4}{256} * 100 = 1.5625\%$ .

e)



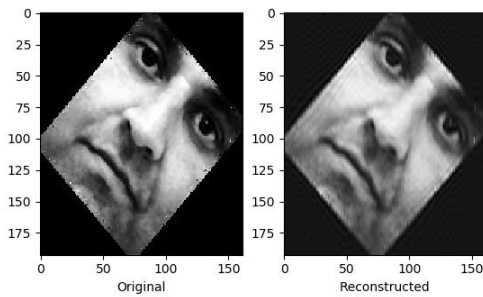
*Figure 6: Reconstruction of the Tesla Cybertruck using all principal components*



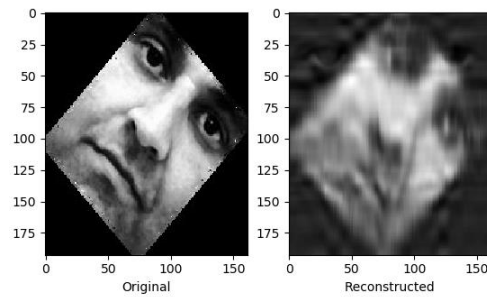
*Figure 7: Reconstruction of the Tesla Cybertruck using 15 principal components*

I tested the ability for the principal components created by using neutral faces to reconstruct a Tesla Cybertruck. The reconstruction in Figure 6 is quite good. Although some of the sharper detail is lost, one can clearly tell that it is a car. Just to see what would happen, I attempted a reconstruction of the truck using only 15 principal components. You can see a blurry car in the background accompanied by a pair of creepy eyes. I'm surprised at how good the reconstruction is when using all the principal components – I was expecting the results in Figure 6 to be about as good as they are in Figure 7.

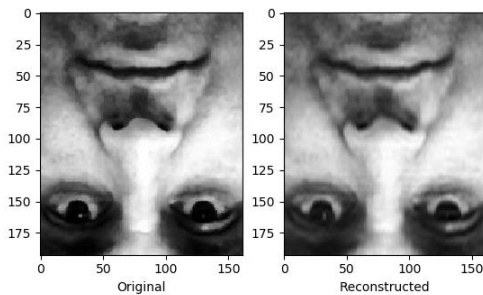
f)



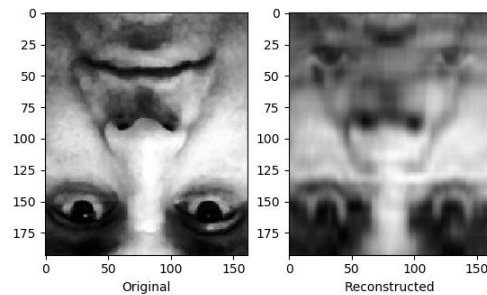
*Figure 8: Reconstruction of neutral face 4a rotated 45 degrees using all principal components*



*Figure 9: Reconstruction of neutral face 4a rotated 45 degrees using 15 principal components*



*Figure 10: Reconstruction of neutral face 4a rotated 180 degrees using all principal components*



*Figure 11: Reconstruction of neutral face 4a rotated 180 degrees using 15 principal components*

I tested the ability for the principal components created from neutral faces to reconstruct image 4a (an image in the set of faces used to create the principal components) after being rotated. In Figures 8 and 10 the recreation is very good – there is just a bit of detail lost in the finer features of the face. When only using 15 principal components the creepy eyes come back, but one can begin to see the eyes and part of the nose of the original image (in the correct orientation). I’m also impressed with the reconstruction results of Figure 8, more so than the reconstruction of Figure 10. I had assumed the greatest variation would be in the vertical direction and the second greatest variation to be in the horizontal direction (simply from observing the organization of features on the faces). Since Figure 10 keeps a vertical and horizontal organization of features I expected the reconstruction to be fine. In Figure 8, however, the

orientation of features is now in a diagonal direction and I expected this to hinder the ability for the principal components to reconstruct the image to the point that Figure 8 would look like Figure 9.

## Appendix

*Code:*

```

1 #####
2 # Joseph Bell #
3 # ECE 269 PCA Eigenfaces Project #
4 # 12/1/2019 #
5 #####
6
7 import matplotlib.pyplot as plt
8 import matplotlib.image as mplimg
9 import os
10 import numpy as np
11 import math
12
13
14 #####
15 # input: absolute path to folder containing face images#
16 # returns: m * L * N array, where m = number of face #
17 # images, L rows in image, N cols in image #
18 #####
19 def vectorize_face_folder(face_folder_path): #converts face images to
matrix
20     face_images = os.listdir(face_folder_path)
21     first_face_path = os.path.join(face_folder_path, face_images[0]) #
Used to initialize np.array to store all faces - assumes all faces are same
size as first image
22     num_of_faces = len(face_images) #
which is true in this case as they've all been cropped
23     first_face = mplimg.imread(first_face_path)
24     faces_vector = np.zeros((num_of_faces, first_face.shape[0],
first_face.shape[1]))
25
26     for i in range(num_of_faces):
27         face_path = os.path.join(face_folder_path, face_images[i])
28         face_read = mplimg.imread(face_path)
29         faces_vector[i,:,:] = face_read
30
31     return faces_vector
32
33 #####
34 # input: m*L*N array, where m = number of face #
35 # images, L rows in image, N cols in image #
36 # returns: L*N array which is the mean of m face #

```

```

37 #             images                                     #
38 #####
39 def calculate_mean_face(faces_vector): #calculates mean of all matrices
that represent faces
40     num_of_faces = faces_vector.shape[0]
41     face_summation = np.zeros(faces_vector[0].shape)
42     for i in range(num_of_faces):
43         face_summation = face_summation + faces_vector[i,:,:]
44
45     mean_face = face_summation/num_of_faces
46
47     return mean_face
48
49 #####
50 # input (faces_vector): m*L*N array, where m = number of face #
51 #             images, L rows in image, N cols in image #
52 # input (mean_face): L*N array of mean face of m face images #
53 # returns: m*L*n array of faces_vector[i,:,:] - mean_face #
54 # #
55 #####
56 def calculate_mean_adjusted_faces(faces_vector, mean_face):
57     mean_adjusted_faces = np.zeros(faces_vector.shape)
58     num_of_faces = faces_vector.shape[0]
59     for i in range(num_of_faces):
60         mean_adjusted_faces[i,:,:] = faces_vector[i,:,:] - mean_face
61
62     return mean_adjusted_faces
63
64 #####
65 # input: m*L*N array, where m = number of face #
66 #             images, L rows in image, N cols in image #
67 # returns: L*L array which is the covariance matrix #
68 #####
69 def calculate_covariance_matrix(mean_adjusted_faces):
70     num_of_faces = mean_adjusted_faces.shape[0]
71     first_face = mean_adjusted_faces[0] #used for initializing
AAT_face_summation shape
72     AAT = np.zeros((first_face.shape[0], first_face.shape[0])) # where A
is a mean adjusted face matrix
73
74     # Used 3D arrays as the computation time is almost identical this
way, however I did read through the quicker method in the paper they used
75     # (they used 2D arrays) and understand it (i.e. why it's quicker and
how it results in the same eigenvalues and scaled eigenvectors.
76     # Also, had already implemented 3D array method in all my code and it
would have been a hassle to change everything
77     # to work with 2D dimensions that also would've been of minor benefit
to computation time (and less intuitive for me - 3D arrays make more
78     # sense to me when dealing with images.
79
80     for i in range(num_of_faces):
81         face = mean_adjusted_faces[i,:,:]
82         AAT = AAT + np.matmul(face, face.T) #if face is 193*162 then
AAT_face_summation is 193*193
83
84     return AAT
85

```

```

86 #####
87 # input: L*L array which is the covariance matrix #
88 # returns: L, L*L array of eigenvalues and eigenvectors #
89 #####
90 def calculate_eigenpairs(covariance_matrix):
91     eigenvalues, eigenvectors = np.linalg.eig(covariance_matrix)
92
93     return eigenvalues, eigenvectors
94
95 #####
96 # input: L array of eigenvalues #
97 # input: L*L array of eigenvectors #
98 # input: num_of_faces is number of PCs for face space #
99 # returns: L*L, L array of eigenvectors and eigenvalues for face space #
100 #####
101 def select_face_space(eigenvalues, eigenvectors, num_of_faces):
102     sorted_eigenvalues = sorted(eigenvalues, reverse=True)
103     face_space_eigenvectors = np.zeros((eigenvectors.shape[0],
num_of_faces)) #each column is an eigenvector
104     face_space_eigenvalues = np.zeros(num_of_faces)
105     for i in range(num_of_faces):
106         eigenvalue = sorted_eigenvalues[i]
107         corresponding_eigenvector_index = np.where(eigenvalues ==
eigenvalue)[0][0]
108
109         corresponding_eigenvector = eigenvectors[:,
corresponding_eigenvector_index]
110
111         face_space_eigenvectors[:,i] = corresponding_eigenvector
112         face_space_eigenvalues[i] = eigenvalue
113
114     return face_space_eigenvectors, face_space_eigenvalues
115
116 #####
117 # input: L array of eigenvalues for face space #
118 # return: no return, just creates bar plot for singular values #
119 #####
120 def create_principal_component_bar_plot(face_space_eigenvalues):
121     number_of_eigenvalues = face_space_eigenvalues.shape[0]
122     singular_values = np.sqrt(face_space_eigenvalues)
123     eigenvalue_index_array = np.linspace(1, number_of_eigenvalues,
number_of_eigenvalues)
124     plt.bar(eigenvalue_index_array, singular_values)
125     plt.xlabel('Singular Value')
126     plt.ylabel('Singular Value Magnitude')
127     plt.show()
128
129 #####
130 # input (face_to_reconstruct): L*N image to reconstruct using face space#
131 # input (face_space_eigenvectors): L*L array of face space eigenvectors #
132 # input (mean_face): L*N mean face (constructed from 190 neutral faces) #
133 # returns: L*N reconstructed image #
134 #####
135 def principal_component_image_reconstruction(face_to_reconstruct,
face_space_eigenvectors, mean_face): #takes in a single image and
reconstructs it by using principal components
136     number_of_faces = face_space_eigenvectors.shape[1]

```



```

137     reconstructed_image = np.zeros(face_to_reconstruct.shape)
138     mean_adjusted_face = face_to_reconstruct - mean_face
139     for i in range(number_of_faces):
140         eigenvector =
face_space_eigenvectors[:,i].reshape((face_space_eigenvectors[:,i].shape[0],1
))
141         w = np.matmul(eigenvector.T, mean_adjusted_face)
142         w_u = np.matmul(eigenvector, w)
143         reconstructed_image = reconstructed_image + w_u
144
145     reconstructed_image = reconstructed_image + mean_face
146     return reconstructed_image
147
148 #####
149 # input (image_reconstruction): L*N reconstructed image #
150 # input (face_to_reconstruct): original image used for reconstruction #
151 # returns: Mean Squared Error of inputs #
152 #####
153 def calculate_MSE(image_reconstruction, face_to_reconstruct):
154     L = image_reconstruction.shape[0]
155     N = image_reconstruction.shape[1]
156
157     MSE = np.sum(np.square(image_reconstruction -
face_to_reconstruct))/(L*N)
158
159     return MSE
160
161 #####
162 # input: L*N rgb image #
163 # returns: L*N gray image #
164 #####
165 def rgb2gray(image):
166     gray_image = np.zeros((image.shape[0], image.shape[1]))
167     gray_image[:, :] = image[:, :, 0]*0.2989 + image[:, :, 1]*0.5870 +
image[:, :, 2]*0.1140
168
169     return gray_image
170
171 #####
172 # input: L*L array of eigenvectors #
173 # returns: L array of cumulative significance used for plotting #
174 #####
175 def calculate_significance_of_eigenvalues(eigenvalues):
176     total_sum = sum(eigenvalues)
177     cumulative_significance = np.zeros(eigenvalues.shape[0])
178     for i in range(eigenvalues.shape[0]):
179         cumulative_significance[i] = sum(eigenvalues[:i])/total_sum
180     return cumulative_significance
181
182 # Folder used to create Principal Components
183 face_folder = 'neutral_faces'
184
185 # Folder used to access images to reconstruct
186 face_folder_reconstruct = 'neutral_faces'
187 #face_folder_reconstruct = 'smiling_faces'
188
189

```

```
190 # Creating absolute paths
191 face_folder_path = os.path.join(os.getcwd(), face_folder)
192 face_folder_path_reconstruct = os.path.join(os.getcwd(),
face_folder_reconstruct)
193
194 ##### Below are all used one at a time for different reconstructions ###
195 #face_to_reconstruct =
mplimg.imread(os.path.join(face_folder_path_reconstruct, '4b.jpg'))
196 #face_to_reconstruct =
mplimg.imread(os.path.join(face_folder_path_reconstruct, '195a.jpg')) #to be
used later to test MSE of image reconstruction
197 #face_to_reconstruct =
rgb2gray(mplimg.imread(os.path.join(os.getcwd(), 'tesla_cybertruck.jpg')))
198 face_to_reconstruct =
mplimg.imread(os.path.join(os.getcwd(), 'neutral_rotate_45.jpg'))
199
200 ##### Code to calculate Principal Components #####
201 faces_vector = vectorize_face_folder(face_folder_path=face_folder_path)
202 mean_face = calculate_mean_face(faces_vector=faces_vector)
203
204 mean_adjusted_faces =
calculate_mean_adjusted_faces(faces_vector=faces_vector, mean_face=mean_face)
205 covariance_matrix =
calculate_covariance_matrix(mean_adjusted_faces=mean_adjusted_faces)
206 eigenvalues, eigenvectors =
calculate_eigenpairs(covariance_matrix=covariance_matrix) #eigenvectors have
euclidean norm of 1
207
208 #99% significance occurs at 54
209 num_of_PC = 54
210 face_space_eigenvectors, face_space_eigenvalues =
select_face_space(eigenvalues=eigenvalues, eigenvectors=eigenvectors,
num_of_faces=num_of_PC) #max value for num_of_faces is total number of
eigenvectors, number is chosen heuristically
211
create_principal_component_bar_plot(face_space_eigenvalues=face_space_eigenva
lues)
212 cumulative_significance =
calculate_significance_of_eigenvalues(eigenvalues=face_space_eigenvalues)
213
#####
#####3
214
215 # Uncomment code to plot cumulative significance chart - only required
once to determine how many principal components to used
216 # for the project. 99% of the sum of magnitudes of all eigenvalues is
obtained at eigenvalue 54. Therefore, 54 principal components
217 # are used.
218
219 indices = np.linspace(1, cumulative_significance.shape[0],
cumulative_significance.shape[0])
220 plt.plot(indices, cumulative_significance)
221 plt.plot(54, 0.99, 'or')
222 plt.xlabel('Principal Component')
223 plt.ylabel('Cumulative Significance')
224 plt.show()
225
```

```
226 # Reconstruct image
227 image_reconstruction =
principal_component_image_reconstruction(face_to_reconstruct=face_to_reconstruct,
face_space_eigenvectors=face_space_eigenvectors, mean_face=mean_face)
228
229 # Plotting original and reconstructed image next to each other
230 f = plt.figure()
231 f.add_subplot(1,2,1)
232 plt.imshow(face_to_reconstruct, cmap='gray')
233 plt.xlabel('Original')
234 f.add_subplot(1,2,2)
235 plt.imshow(image_reconstruction, cmap='gray')
236 plt.xlabel('Reconstructed')
237 plt.show()
238
239 ##### Uncomment code to calculate MSE and create plots of MSE vs number
of principal components #####
240
241 MSE = calculate_MSE(image_reconstruction=image_reconstruction,
face_to_reconstruct=face_to_reconstruct)
242 PC_indices = np.linspace(1,num_of_PC,num_of_PC)
243 MSE_array = np.zeros(num_of_PC)
244
245 for i in range(num_of_PC):
246     face_space_eigenvectors, face_space_eigenvalues =
select_face_space(eigenvalues=eigenvalues, eigenvectors=eigenvectors,
num_of_faces=i)
247     image_reconstruction =
principal_component_image_reconstruction(face_to_reconstruct=face_to_reconstruct,
face_space_eigenvectors=face_space_eigenvectors, mean_face=mean_face)
248     MSE = calculate_MSE(image_reconstruction=image_reconstruction,
face_to_reconstruct=face_to_reconstruct)
249     MSE_array[i] = MSE
250
251 plt.plot(PC_indices, np.sqrt(MSE_array))
252 plt.xlabel('Number of Principal Components')
253 plt.ylabel('Square root of MSE')
254 plt.show()
```